

# Project 9 — Show and Tell

Jonathan Sumner Evans (jonathanevans@mines.edu)

Sam Sartor (ssartor@mines.edu)

Daichi Jameson (djameson@mines.edu)

May 4, 2017

## 1 Application Overview

For our final project, we decided to build a web application for Mines students to show off their cool projects on the ACM TV in Brown Building West. We call the application *Show and Tell* after the age-old first-grade tradition. The core functionality *Show and Tell* is that users can create teams and upload projects. These projects are then verified by a system administrator, and once verified, the projects are ready for display on the ACM TV. Displaying will be handled by a separate program, Visplay, running on the TV, which has yet to be implemented.

### 1.1 Application Schema

The core entities in our application are users, teams, and projects. In addition to these entities, *Show and Tell* tracks user sessions and project assets. The following ERD (Figure 1) describes the relationships between these entities.

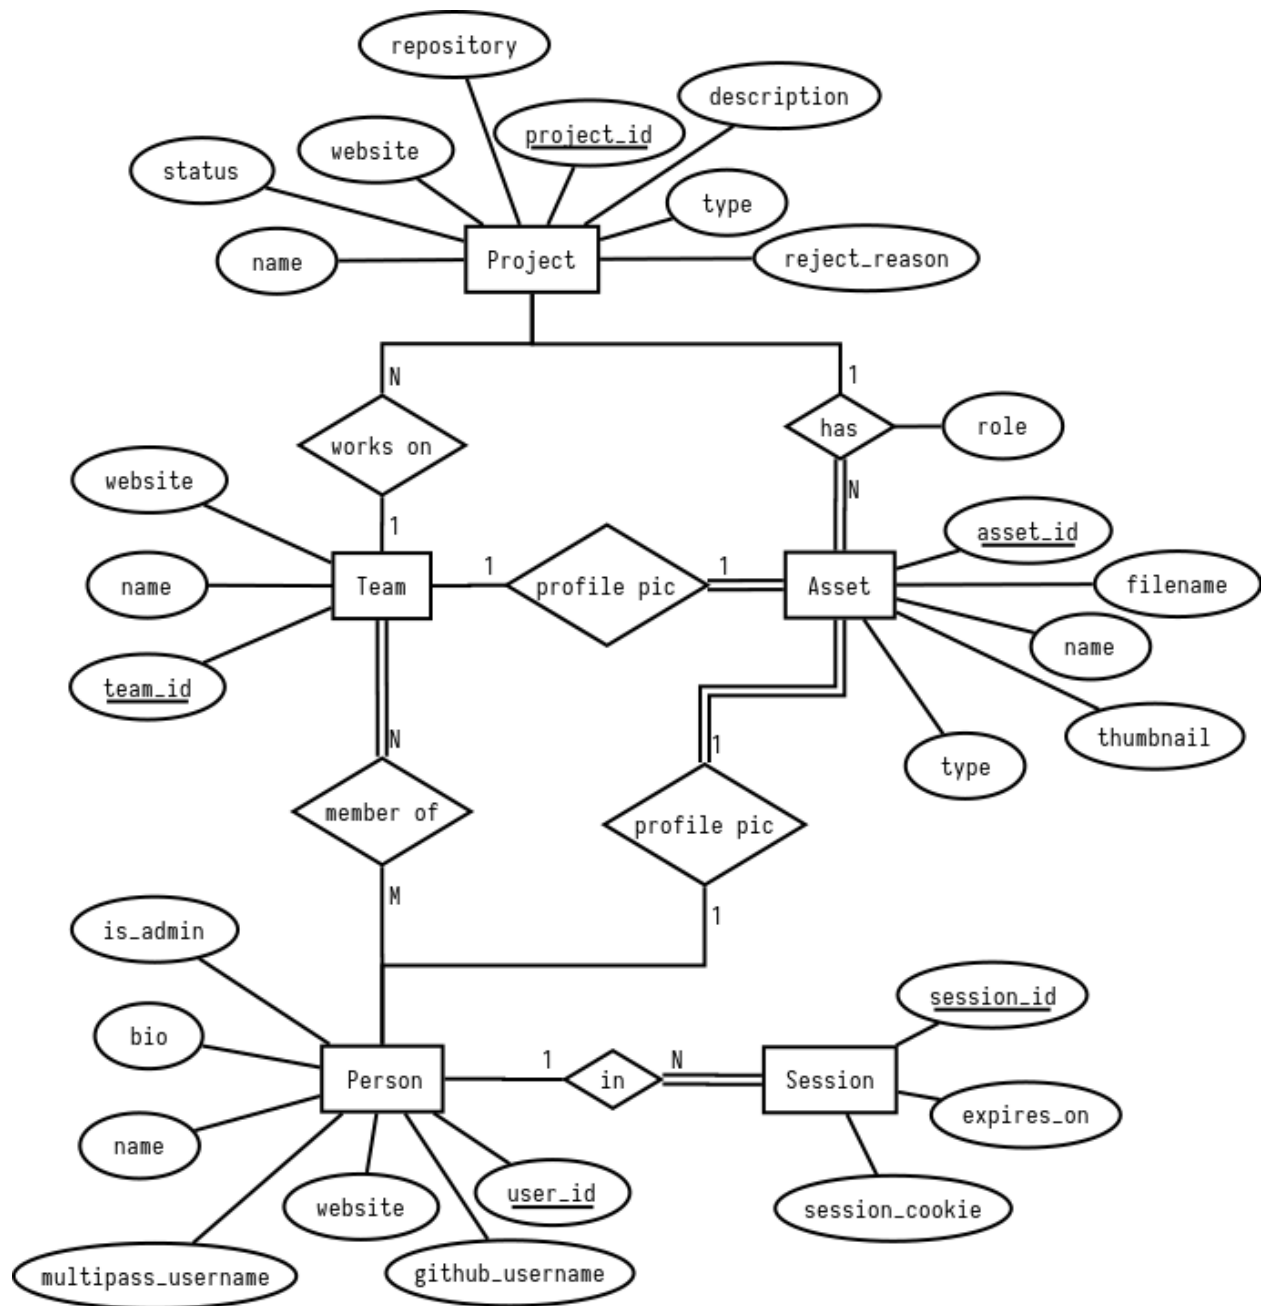


Figure 1: The ERD for *Show and Tell*

## 1.2 Schema Explanation

### 1.2.1 Entities

The database contains the following entities: Projects, Teams, Assets, People, and Sessions. The Project, Team and Person entities are self explanatory. The Asset entity represents a single uploaded asset. These assets could be profile pictures or project assets. Project assets can be any type of file such as program executables, videos, or source code to name a few. The Session entity represents a user session and associates a user of *Show and Tell* in their browser with their user account in the database.

### 1.2.2 Relationship

People (users) can be in multiple Sessions. This relationship links browser sessions with users in the database. People can be members of as many of teams as they would like and teams may have multiple team members (m-n relationship). Teams can work on multiple projects, but only one team can work on any given project (1-n relationship). Each Project can have multiple assets and each asset associated with a project has a role which could be anything from project executable to readme file. Additionally, every user and team can have a single profile picture asset. Assets can have at most one owner which can be either a person, team, or project.

The reasons for most of the relationships are fairly intuitive, but we made a few interesting design decisions. First, all assets are stored in the same table. We chose this design because we wanted to have a single registry of files in the system rather than having multiple places that we need to manage files. A second interesting design choice was to add a role to the Project-has-Asset relationship. We chose to do this because in the future, we want to store project metadata such as configuration files or Dockerfiles in the system for when we display the project on the ACM TV.

It would be better if each asset has a reference to one, and only one, owner (which is either a person, project, or team). Then the role of that asset (profile pic, readme, slide, etc.) determines its actual role. Getting a profile pic would then be done using the query

---

```
1 SELECT * FROM assets WHERE owner=person_id AND role='profile_pic'
```

---

rather than

---

```
1 SELECT * FROM people AS p JOIN assets AS a ON a.id=p.profile_pic WHERE p.id=person
```

---

Unfortunately, this relationship is not easy to represent in SQLAlchemy and as we were developing the models for the application we did not fully understand the framework and we haven't gotten around to fixing it.

## 2 Application Implementation

### 2.1 Authentication and Authorization

All actions that require authentication perform a lookup of the session id cookie in the Session table to determine if the user is logged in. We then use the user to determine what projects and teams the user has permission to edit.

### 2.2 Application Structure

We used an Model-View-Controller (MVC) architectural pattern for our application. This pattern is very common among modern web applications and was especially suited for our application.

#### 2.2.1 Model

The models in our application are essentially wrappers around the SQL entities in the database. To communicate with the database we used SQLAlchemy, an Object Relational Mapping (ORM) tool that can build and manipulate databases. The primary use of the database was to aggregate projects and project information, which would then be forwarded onto the ACM TV after approval. One of the major design decisions made was to store the data for the project assets on the file system of the server and only keep basic information about the asset, such as the asset name, in the SQL database. The model logic handled the linkage between the filesystem and the database. This allowed for the database to be significantly lighter, and also allowed us to circumvent some restrictions that that SQL databases

have with directly storing files, such as size limits. This can also give a performance increase by reducing the number of copies necessary to serve an asset (Bottle can stream directly from disk rather than copying through the ORM).

### **2.2.2 View**

We chose to try and minimize the amount of JavaScript required in our application. To accomplish this, we tried to maximize the amount of HTML dynamically generated on the server. This required a templating engine and we chose to use Kajiki. Kajiki dynamically generates HTML using data from the application including directly from our SQL database. This is done through template files which are largely pure XHTML, but can include Python within the XHTML elements attributes. This embedded Python can filter, repeat, and alter its associated HTML elements. One advantage of this method is that we were able to generate large and complex XHTML without much code overhead.

### **2.2.3 Controller**

We used Bottle to route HTTP requests to the corresponding controllers. Bottle is a fast, simple and lightweight WSGI micro web-framework for Python and is the driving force of our application. When an HTTP request comes into the server, Bottle calls the corresponding Python function to construct a response to return to the user. Bottle allows us to determine what web page needs to be served to a user, grabs the necessary database information from the SQLAlchemy models, and then passes the information to Kajiki to interact with the user. This library is the glue that connects all of the pieces together.

## **2.3 Additional Libraries**

In addition to the Python libraries mentioned above, we used a few libraries to build the UI. The primary libraries that we used were Bootstrap, jQuery, and Compass. Bootstrap provides a large variety of nice styles for HTML components and provides some convenient JavaScript utilities. jQuery gives us convenient DOM manipulation abstractions in JavaScript. We used Compass, a CSS compiler, which provides some nice features such as file importing and easy minification of the resulting CSS.

### 3 Implementation Difficulties

Making sure asset information stayed consistent between different parts of the ORM and the SQL database proved to be a challenging task. While the ORM easily sets up and builds the necessary SQL tables, additional configuration is required to ensure that calls to remove assets are propagated throughout the ORM and the database. One specific issue that we encountered during development was that the project object created in the ORM did not recognize when rows were dropped from the asset and asset cross-reference table. By adding in a couple configurations, we were able to remove assets directly from the project, which would then drop the corresponding rows in the project-asset cross reference table and the asset table. Once we figured out how to use this feature, the ORM made it extremely intuitive to work with.

Another issue we ran across was the orphaning of assets in the file system. While the internal issues between the SQL database and the ORM were solved with some additional ORM configuration, asset data still remained on the file system after the asset information was deleted off of the ORM and database. The orphaning issue was solved by using the Python standard library to delete the file data off of the file system before removing the rows off of the table.

### 4 What We Learned

None of us had used an ORM before this project. While there were a couple of exceptions, setting up the database schema was very simple and intuitive. Each table is simply a new Python class. Columns and relationships can be added to a table by adding variables of type Column and Relationship, respectively. Manipulating and querying variables was also very convenient since queries were returned as Python objects. This allowed the SQL data to work very well with the other Python libraries, and allowed us to pass information from queries directly from the ORM to the XHTML templates to display user, project, and team information.

Another valuable skill we learned was integrating a file system with the SQL database in order to store files. With the help of the Python standard library, storing files on disk was pretty straightforward. After we retrieved the file using Bottle, we took the file information (such as the file name) and added it to the asset table on the SQL server. The file data was

then assigned a UUID and stored on the file system. Whenever we needed to serve a file, we would use Bottle to take the data off of the disk, and then attach the name of the file from the SQL server onto the data before serving it to the user.

We also learned a great deal about web application design. Although we were familiar with the MVC paradigm, we had never implemented a full MVC app before. Additionally, we used a variety of libraries that we had never used before or had had minimal experience with. Although none of the libraries had an especially steep learning curve, learning all of the libraries at the same time was not trivial.

Because this was a full-featured web application, we had to have a project management structure. We used Trello to track todos and bugs and although Trello is easy to use, effectively using it to manage the project proved to be non-trivial. Our initial todos were too broad so as we went through the development process, we split them into many tasks.

## 5 Conclusion

Through the process of building this app, we learned about how to integrate a database into a web application in a practical, real-world application. Through this process, we learned about creating application architecture, we learned a variety of new libraries, and we learned about how to manage developing a large-scale application. Overall, the result of our efforts were very successful.

## 6 Notes on the Submission

We have put the application up on AWS. It can be accessed via <http://showandtell.samsartor.com> (it redirects to the nasty AWS site).

The source code is also available on GitHub here:

<https://github.com/ColoradoSchoolOfMines/show-and-tell>

We have also included a zip file with the source code. The database schema is also on Jonathan Sumner Evans' database schema on flowers.